

Principles of Robot Autonomy I

Multi-sensor perception and sensor fusion



Stanford
University

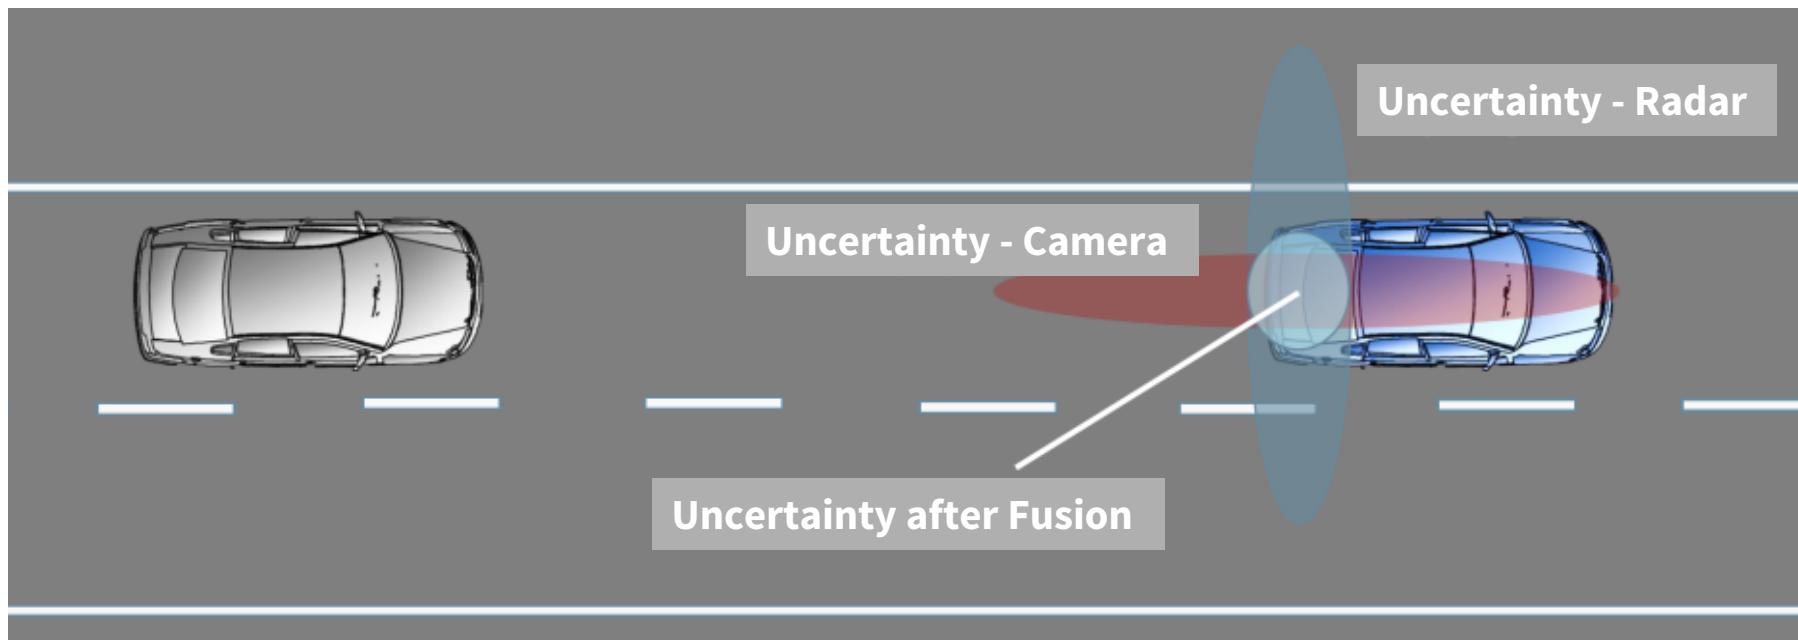


Today's lecture

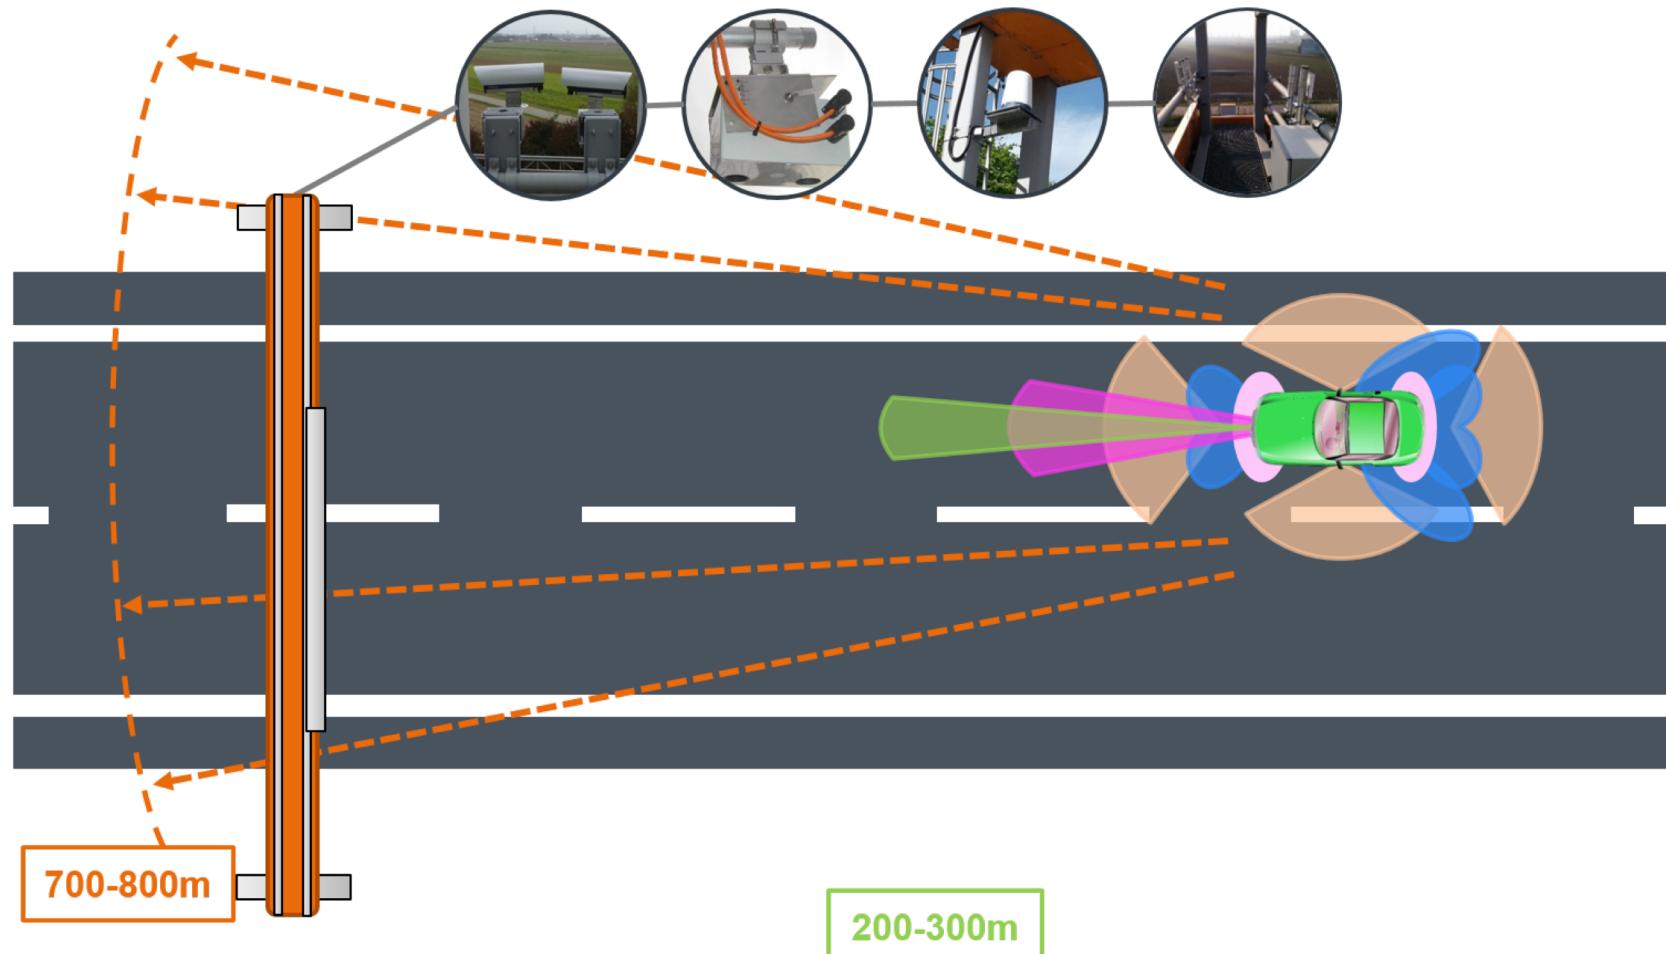
- Aim
 - Introduce the topic of multi-sensor perception and sensor fusion
 - Learn about Kalman filtering applied to sensor fusion
 - Devise a sensor fusion algorithm for position estimation
- Readings
 - F. Gustafsson. Statistical Sensor Fusion. 2010.
 - D. Simon. Optimal State Estimation: Kalman, H_∞ , and Nonlinear Approaches. 2006.

Multi-sensor perception

- Uncertainty reduction



Using stationary sensors



Single-sensor vs multi-sensor perception

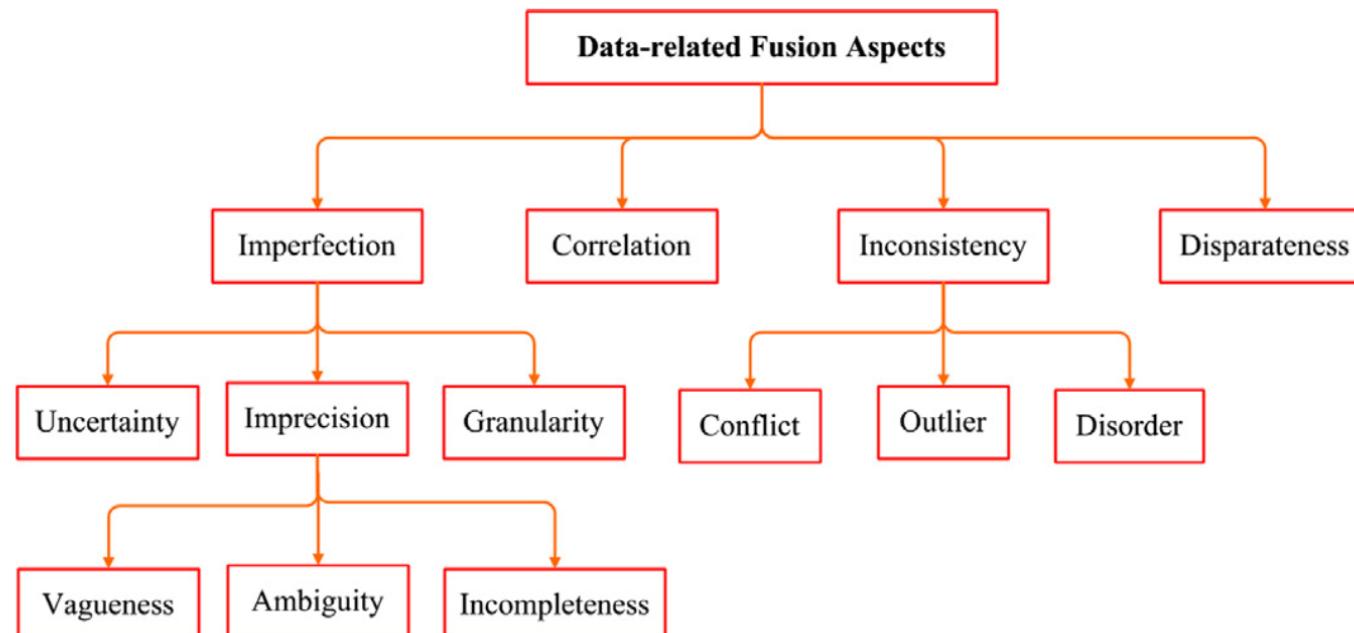
- Drawbacks of single-sensor perception
 - Limited range and field of view
 - Performance is susceptible to common environmental conditions
 - Range determination is not as accurate as required
 - Detection of artefacts, so-called false positives
- Multi-sensor perception might compensate these, and provide:
 - Increased classification accuracy of objects
 - Improved state estimation accuracy
 - Improved robustness for instance in adverse weather conditions
 - Increased availability
 - Enlarged field of view

Sensor fusion taxonomies

- Data-related taxonomy
- Fusion level taxonomy
- Fusion classes taxonomy
- Architectural taxonomy

Data-related taxonomy

- The most interesting data-related fusion aspect is the inherent imperfection of the sensory data
- The data-related taxonomy provides us with a checklist of underlying data issues and how to deal with them

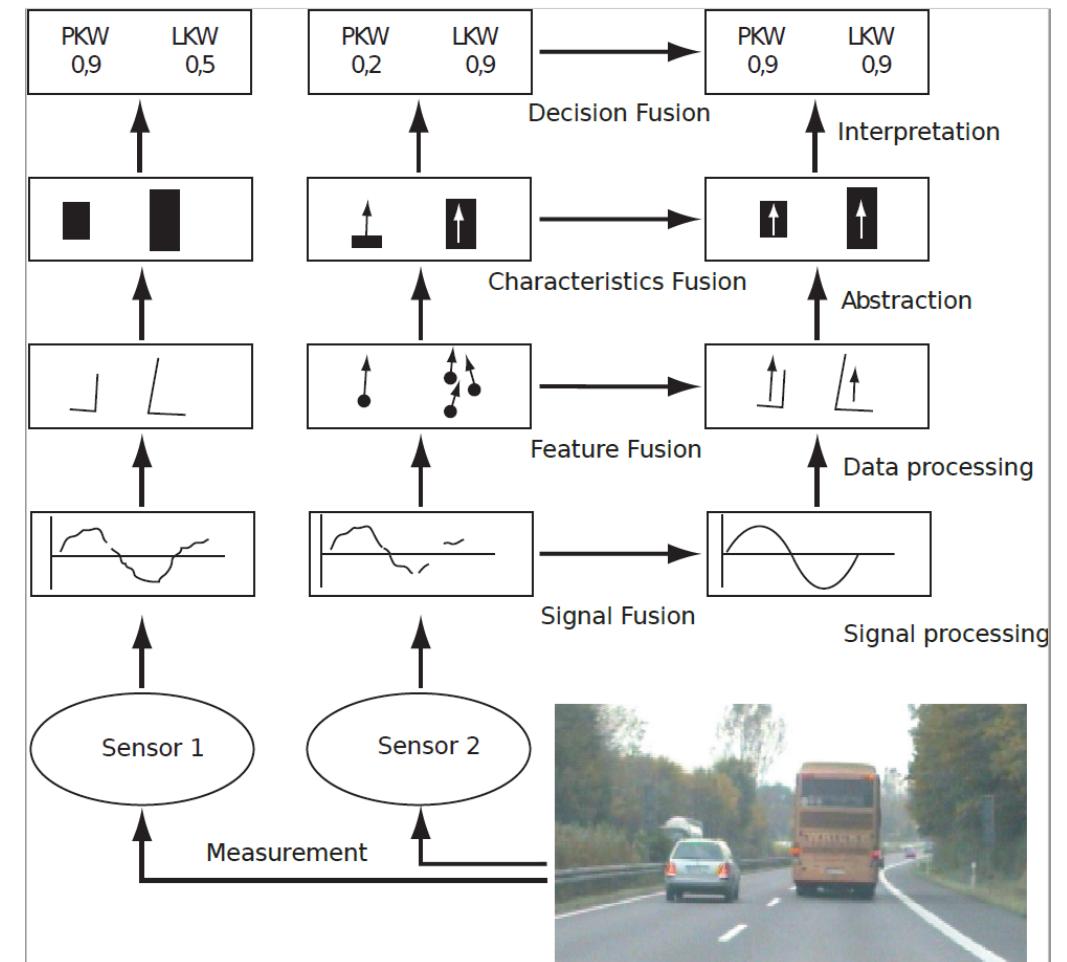


Data-related taxonomy

- Sensory data makes a statement about the environment
 - "The distance to the nearest car is 35.12 m"
- Due to the inherent data imprecision, we have to deal with:
 - **Uncertainty:** The distance to the nearest car is more than 20 m with 80% probability
 - **Vagueness:** The distance to the nearest car is more than 20 m with 80% probability, and we are 90% confident in this statement
 - **Ambiguity**
 - **Incompleteness**
- The underlying data can contain multiple imperfections at once

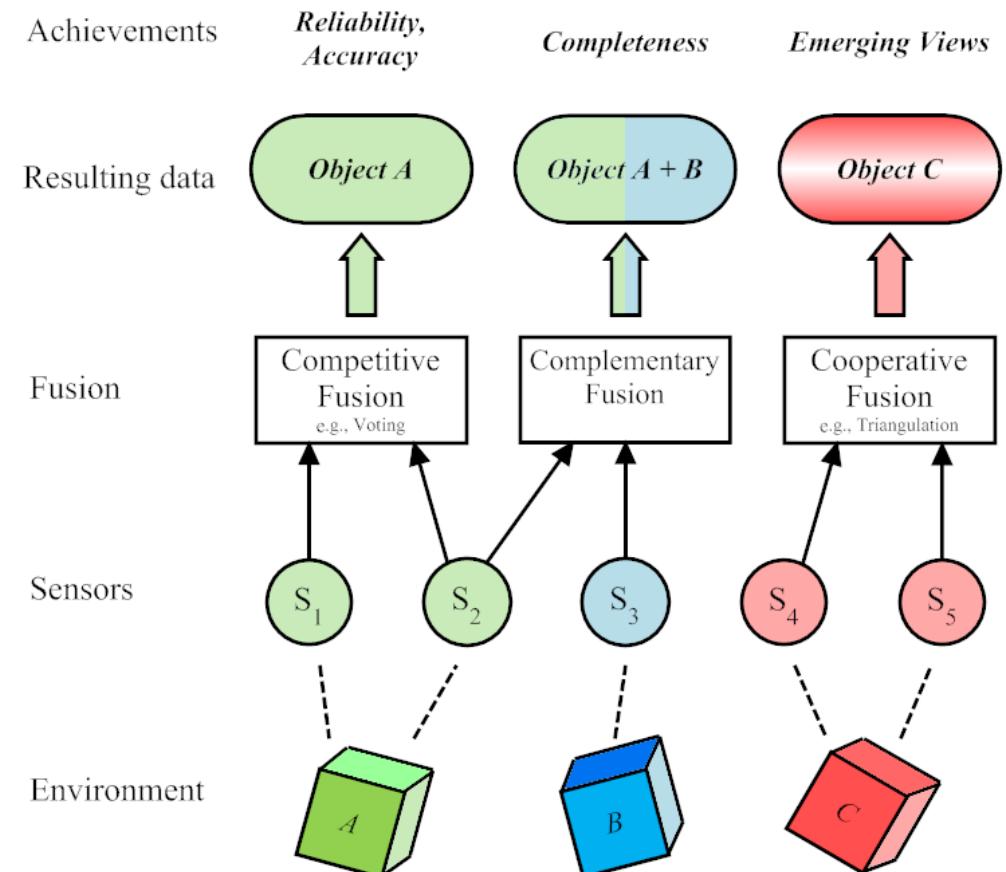
Fusion level taxonomy

- Fusion is typically divided into three levels of abstraction:
 - Low-level fusion
 - Intermediate-level fusion
 - High-level fusion
- They respectively fuse:
 - Signals
 - Features and characteristics
 - Decisions



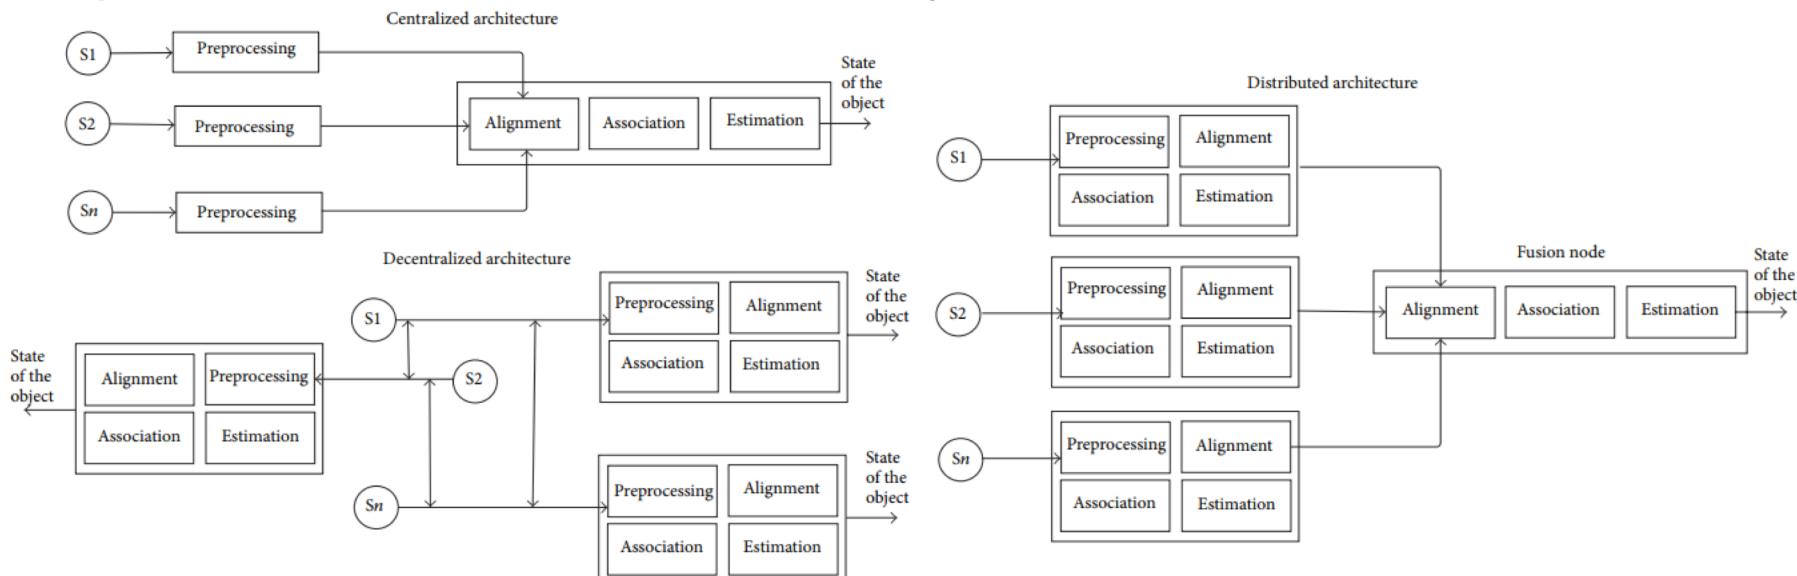
Fusion class taxonomy

- Competitive fusion
 - is used when redundant sensors measure the same quantity, in order to reduce the overall uncertainty
- Complementary fusion
 - is used when sensors provide a complementary information about the environment, for instance distance sensors with different ranges
- Cooperative fusion
 - is used when the required information can not be inferred from a single sensor (e.g. GPS localization and stereo vision)



Architectural taxonomy

- The **centralized** architecture is theoretically optimal, but scales badly with respect to communication and processing
- The **decentralized** architecture is a collection of autonomous centralized systems, and has the same scaling issues
- The **distributed** architecture scales better, but can lead to information loss because each sensor processes its information locally



Bayesian statistics in multi-sensor data fusion

- **Basic premise:** all unknowns are treated as random variables and the knowledge of these quantities is summarized via a probability distribution
 - This includes the observed data, any missing data, noise, unknown parameters, and models
- Bayesian statistics provides
 - a framework for **quantifying objective and subjective uncertainties**
 - principled methods for **model estimation and comparison** and the **classification of new observations**
 - a **natural way to combine different sensor observations**
 - principle methods for dealing **with missing information**

Sensor fusion – a simple example

- **Problem:** determine the distance to n objects using measurements from two sensors
- Assumptions:
 - Both sensors have the same field of view
 - First sensor has a higher precision than the second sensor
 - Consider the simplest case ($n=1$)
- How to fuse these measurements properly?

Sensor fusion – a simple example

- Sensors provide redundant measurements of the same physical quantity (distance)
- To incorporate the precision information → measurements are assumed to be **normally distributed random variables**
- Specifically, the univariate Gaussian distributions are:

$$d_1(x) = (2\pi\sigma_1^2)^{-\frac{1}{2}} \exp\left(-\frac{1}{2}\frac{(x - \mu_1)^2}{\sigma_1^2}\right) \sim \mathcal{N}(\mu_1, \sigma_1^2)$$

$$d_2(x) = (2\pi\sigma_2^2)^{-\frac{1}{2}} \exp\left(-\frac{1}{2}\frac{(x - \mu_2)^2}{\sigma_2^2}\right) \sim \mathcal{N}(\mu_2, \sigma_2^2)$$

Sensor fusion – a simple example

- Assumption from before:
 - First sensor has a higher precision than the second sensor
- This can be captured as: $\sigma_1^2 < \sigma_2^2$
- Problem is to find $d(x) \sim \mathcal{N}(\mu, \sigma^2)$
- The idea is to combine the previous Gaussian distributions

$$d(x) = d_1(x) \cdot d_2(x) = (4\pi^2 \sigma_1^2 \sigma_2^2)^{-\frac{1}{2}} \exp \left(-\frac{1}{2} \left(\frac{(x - \mu_1)^2}{\sigma_1^2} + \frac{(x - \mu_2)^2}{\sigma_2^2} \right) \right)$$

Sensor fusion – a simple example

- Re-arranging the expression in the exponent and dividing the numerator and denominator by $(\sigma_1^2 + \sigma_2^2)$ give

$$\begin{aligned} -\frac{1}{2} \left(\frac{(x - \mu_1)^2}{\sigma_1^2} + \frac{(x - \mu_2)^2}{\sigma_2^2} \right) &= -\frac{1}{2} \frac{(\sigma_1^2 + \sigma_2^2)x^2 - 2(\sigma_2^2\mu_1 + \sigma_1^2\mu_2)x + (\sigma_2^2\mu_1^2 + \sigma_1^2\mu_2^2)}{\sigma_1^2\sigma_2^2} \\ &= -\frac{1}{2} \frac{x^2 - 2\frac{\mu_1\sigma_2^2 + \mu_2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}x + \frac{\mu_1^2\sigma_2^2 + \mu_2^2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}}{\frac{\sigma_1^2\sigma_2^2}{\sigma_1^2 + \sigma_2^2}} \end{aligned}$$

- To obtain an expression of form $x^2 - 2\mu x + \mu^2 = (x - \mu)^2$ in the numerator, it is necessary to add and subtract the square of the middle term

Sensor fusion – a simple example

$$-\frac{1}{2} \frac{x^2 - 2\frac{\mu_1\sigma_2^2 + \mu_2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}x + \left(\frac{\mu_1\sigma_2^2 + \mu_2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}\right)^2 - \left(\frac{\mu_1\sigma_2^2 + \mu_2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}\right)^2 + \frac{\mu_1^2\sigma_2^2 + \mu_2^2\sigma_1^2}{\sigma_1^2 + \sigma_2^2}}{\frac{\sigma_1^2\sigma_2^2}{\sigma_1^2 + \sigma_2^2}}$$

- The expression in the exponent becomes

$$-\frac{1}{2} \frac{(x - \mu)^2 - \mu^2 + s}{\sigma^2} = -\frac{1}{2} \frac{(x - \mu)^2}{\sigma^2} + \frac{\mu^2 - s}{2\sigma^2}$$

Sensor fusion – a simple example

- Putting everything together leads to the final distribution which represents the fused information

$$\begin{aligned}d(x) &= (2\pi\sigma_1\sigma_2)^{-1} \exp\left(-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2} + \frac{\mu^2-s}{2\sigma^2}\right) \\&= (2\pi\sigma_1\sigma_2)^{-1} \exp\left(\frac{\mu^2-s}{2\sigma^2}\right) \exp\left(-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}\right) \\&= C \cdot \exp\left(-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}\right)\end{aligned}$$

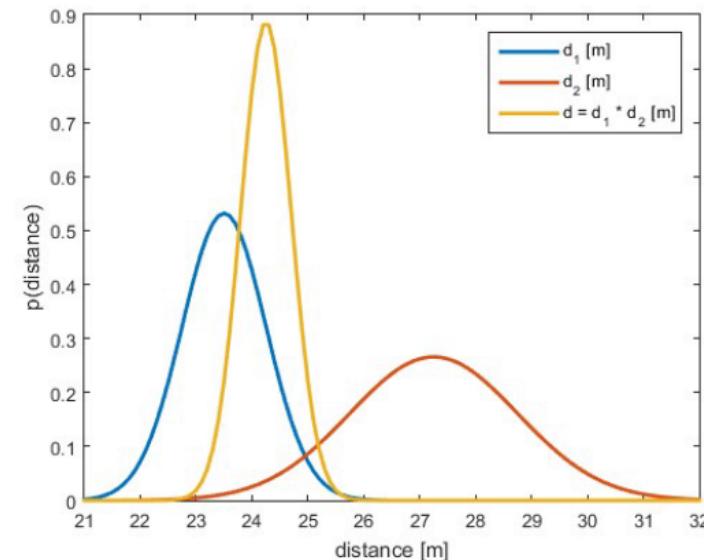
Sensor fusion – a simple example

- Mean value and variance are

$$\mu = \frac{\mu_1 \sigma_2^2 + \mu_2 \sigma_1^2}{\sigma_1^2 + \sigma_2^2}$$

$$\sigma = \frac{\sigma_1^2 + \sigma_2^2}{\sigma_1^2 \cdot \sigma_2^2}$$

- The fused value is the **weighted average** of the measurements
- The **weighting** favors the sensor with higher precision
- The overall **uncertainty decreases**



Kalman filter (KF) – again

- Assumption #1: linear dynamics

$$x_t = A_t x_{t-1} + B_t u_t + \epsilon_t$$

- i.i.d .process noise ϵ_t is $\mathcal{N}(0, R_t)$
- Assumption #1 implies that the probabilistic generative model is Gaussian

$$p(x_t | u_t, x_{t-1}) = \det(2\pi R_t)^{-\frac{1}{2}} \exp \left(-\frac{1}{2} (x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t) \right)$$

Kalman filter (KF)

- Assumption #2: linear measurement model

$$z_t = C_t x_t + \delta_t$$

- i.i.d. measurement noise δ_t is $\mathcal{N}(0, Q_t)$
- Assumption #2 implies that the measurement probability is Gaussian

$$p(z_t | x_t) = \det(2\pi Q_t)^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(z_t - C_t x_t)^T Q_t^{-1} (z_t - C_t x_t)\right)$$

Kalman filter (KF)

- Assumption #3: the initial belief is Gaussian

$$bel(x_0) = p(x_0) = \det(2\pi\Sigma_0)^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(x_0 - \mu_0)^T \Sigma_0^{-1} (x_0 - \mu_0)\right)$$

- **Key fact:** These three assumptions ensure that the posterior $bel(x_t)$ is Gaussian for all t , i.e., $bel(x_t) = \mathcal{N}(\mu_t, \Sigma_t)$
- Note:
 - KF implements a belief computation for continuous states
 - Gaussians are unimodal → commitment to single-hypothesis filtering

Kalman filter: algorithm revisited

Prediction

Project state ahead

$$\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$$

Project covariance ahead

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$$

Correction

Compute Kalman gain

$$K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$$

Update estimate with new measurement

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

Update covariance

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

Prediction:
 $\overline{bel}(x_t)$

Correction:
 $bel(x_t)$

$bel(x_{t-1})$

Data: $(\mu_{t-1}, \Sigma_{t-1}), u_t, z_t$
Result: (μ_t, Σ_t)

$$\left\{ \begin{array}{l} \bar{\mu}_t = A_t \mu_{t-1} + B_t u_t ; \\ \bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t; \end{array} \right.$$

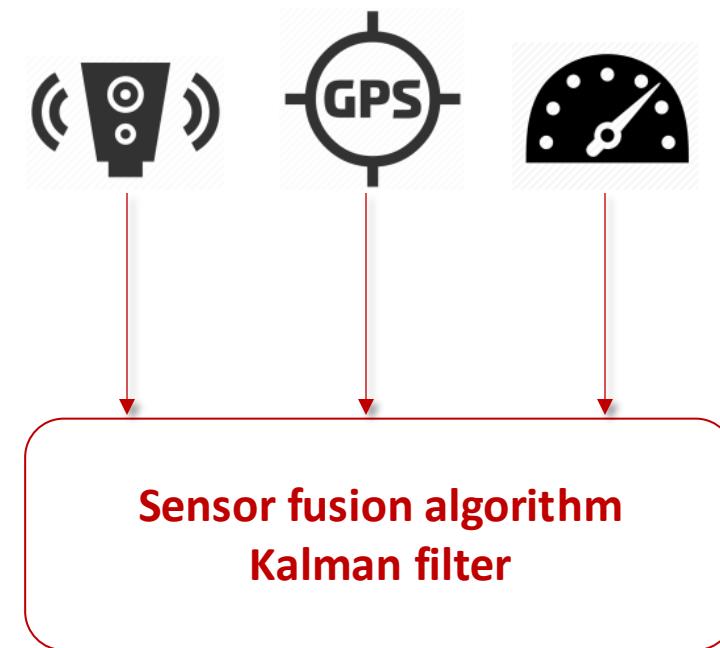
$$\left\{ \begin{array}{l} K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}; \\ \mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t); \\ \Sigma_t = (I - K_t C_t) \bar{\Sigma}_t; \end{array} \right.$$

Return (μ_t, Σ_t)

$bel(x_t)$

Sensor fusion example

- **Problem:** Estimate position, velocity, and acceleration of a vehicle from noisy position and acceleration measurements
- Assumptions:
 - Single track model for the vehicle
 - Lidar provides position measurements with low precision
 - GPS provides position measurements with high precision
 - IMU provides acceleration measurements
- Sensor fusion is done using the **Kalman filter**



Sensor fusion example: Motion model

- **State vector:** $\mu_t = [p \quad v \quad a]^T$
- Change of the state over time is captured by the **motion model**

$$p_t = p_{t-1} + T_s v_{t-1} + \frac{T_s^2}{2} a_{t-1} + \epsilon_{pt}$$

$$v_t = v_{t-1} + T_s a_{t-1} + \epsilon_{vt}$$

$$a_t = a_{t-1} + \epsilon_{at}$$

- T_s represents sampling time

Sensor fusion example: Motion model

- The motion model can be represented in matrix form

$$\begin{bmatrix} p \\ v \\ a \end{bmatrix}_t = \underbrace{\begin{bmatrix} 1 & T_s & \frac{T_s^2}{2} \\ 0 & 1 & T_s \\ 0 & 0 & 1 \end{bmatrix}}_{\text{State transition matrix}} \begin{bmatrix} p \\ v \\ a \end{bmatrix}_{t-1} + \underbrace{\begin{bmatrix} \epsilon_p \\ \epsilon_v \\ \epsilon_a \end{bmatrix}}_{\text{Process noise}}_t$$

State vector *State transition matrix* *Process noise*

$$\mu = A_t \mu_{t-1} + \epsilon_t$$

where ϵ_t is independent process noise distributed as $\mathcal{N}(0, R_t)$

Sensor fusion example: Measurement model

- The **measurement model** defines a mapping from the state space to the measurement space
- For this example, two possible fusion scenarios will be considered:
 1. Lidar+IMU
 2. Lidar+GPS+IMU
- In the first scenario, only measurements from Lidar and IMU are available
 - Assumption: Lidar provides low precision measurement (noisy data)
- In the second scenario, high precision GPS measurements are also available

Sensor fusion example: Measurement model

- First scenario – measurement model is given

$$\begin{bmatrix} p_{lidar} \\ a_{imu} \end{bmatrix}_t = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p \\ v \\ a \end{bmatrix}_t + \begin{bmatrix} \delta_{lidar} \\ \delta_{imu} \end{bmatrix}_t$$

Measurement vector *Measurement matrix* *State vector* *Measurement noise*

$$z_t = C_t \mu_t + \delta_t$$

where δ_t is independent measurement noise distributed as $\mathcal{N}(0, Q_t)$

Sensor fusion example: Initialization

- Choosing the **initial state vector μ_0** - depends on available information

- If there is a-priori knowledge – initialization is done with known values
- If there is a lack of information – initial state is chosen to be zero
- For this example the initial state vector is set to zero

$$\mu_0 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

- Choosing the **initial covariance matrix Σ_0** - should be defined based on the initialization error

- If the initial state is not very close to the correct state - Σ_0 will have large values
- If the initial state is close to the correct state - Σ_0 will have small values

$$\Sigma_0 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

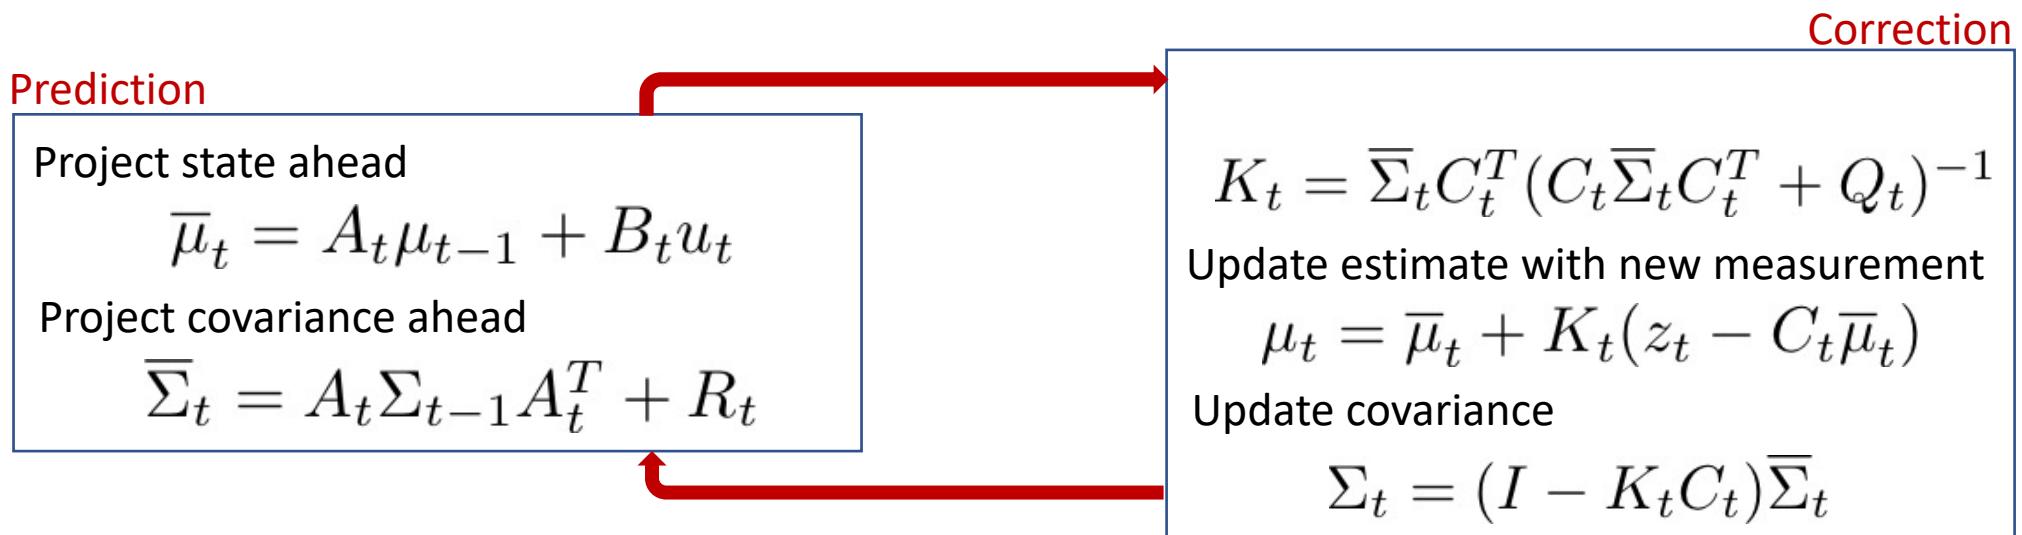
Sensor fusion example: Noise model tuning

- The **process noise covariance matrix R_t** - describes the confidence in the system model
 - Small values indicate higher confidence – predicted values are more weighted
 - Large values indicate lower confidence – measurements become dominant
- The **measurement noise covariance matrix Q_t** - describes the confidence in the measurements
 - Has a similar interpretation as R_t
 - Both matrices need to be symmetric and positive definite

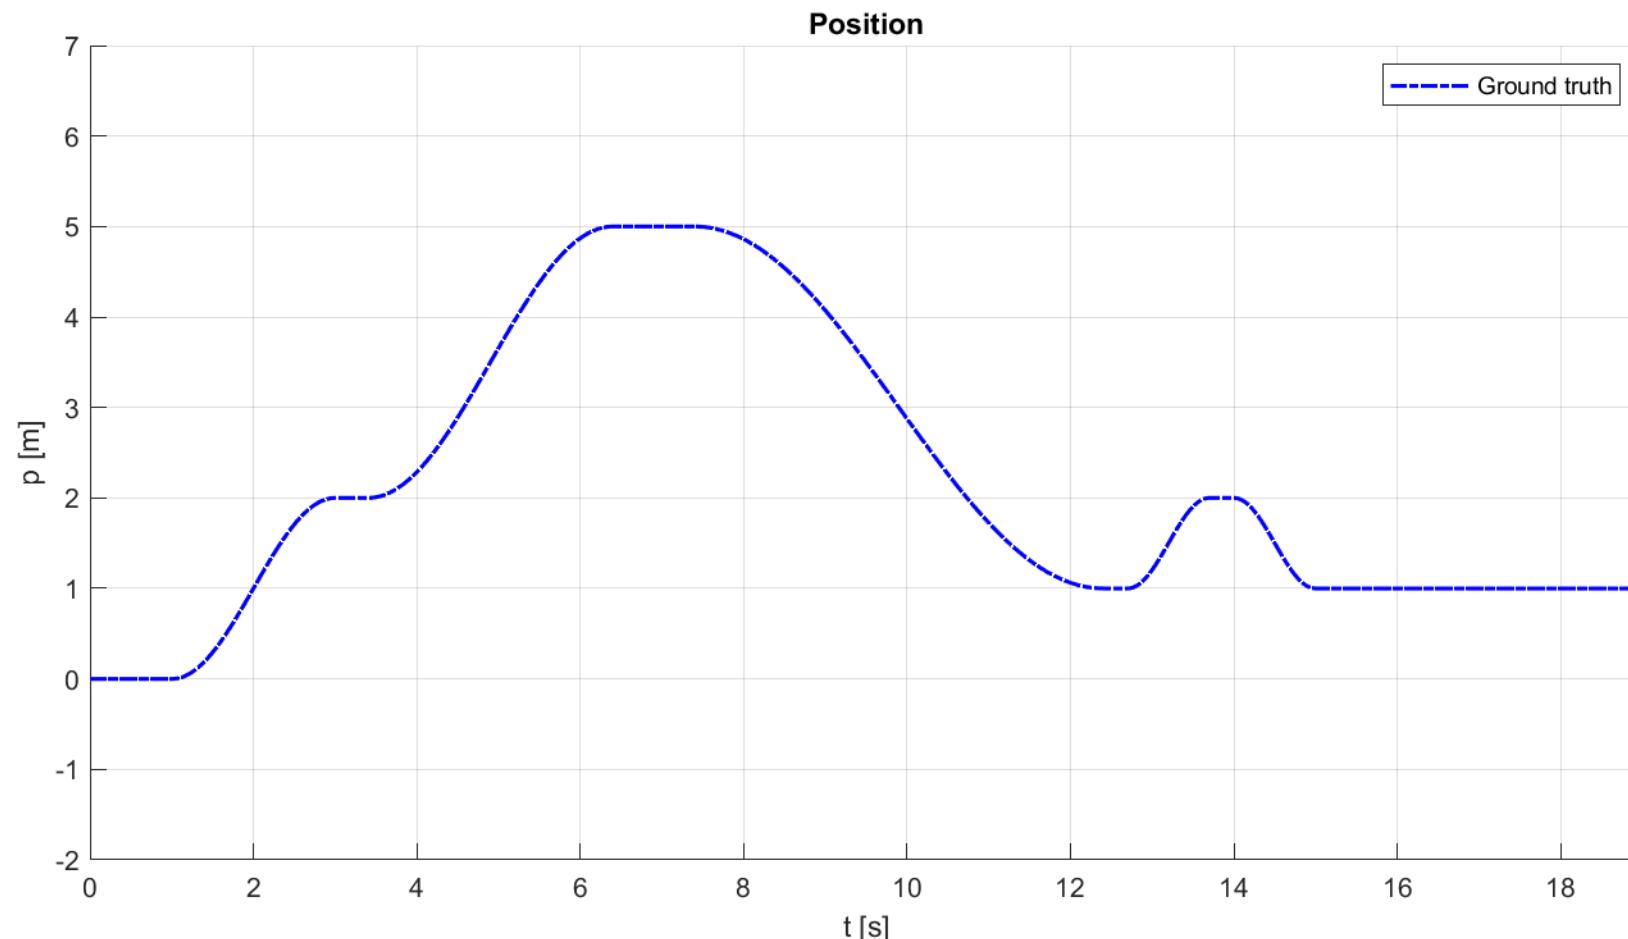
$$R_t = \begin{bmatrix} 0.05 & 0 & 0 \\ 0 & 0.001 & 0 \\ 0 & 0 & 0.05 \end{bmatrix} \quad Q_t = \begin{bmatrix} \sigma_{lidar}^2 & 0 \\ 0 & \sigma_{imu}^2 \end{bmatrix}$$

Sensor fusion example: Algorithm

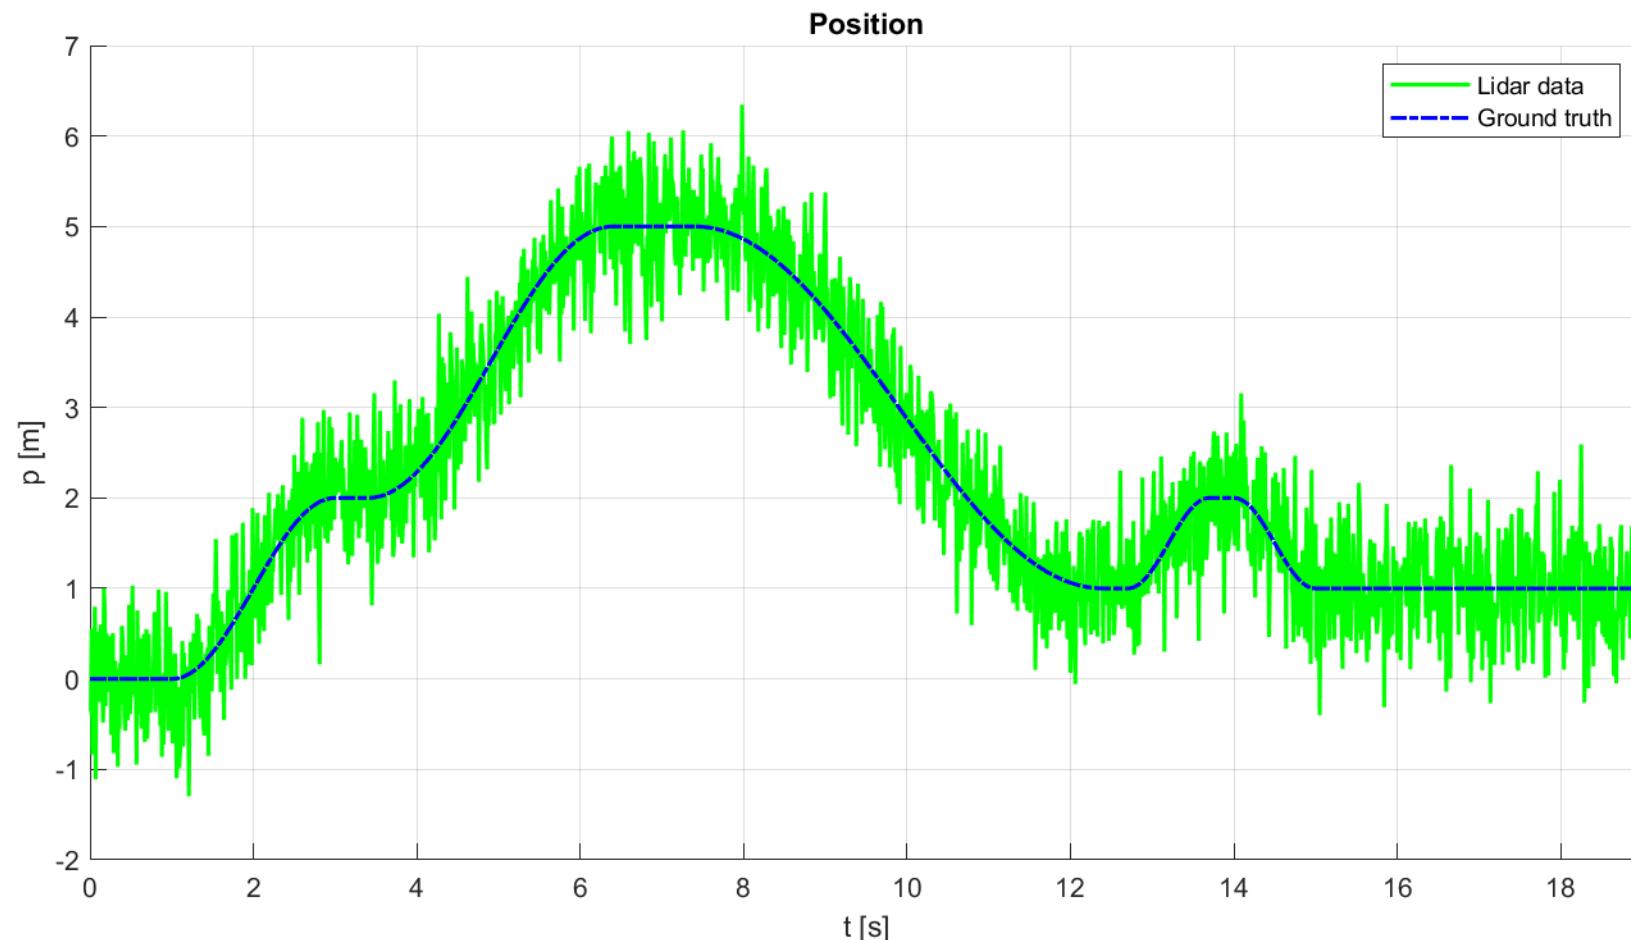
- Estimation results are obtained using the prediction-correction scheme



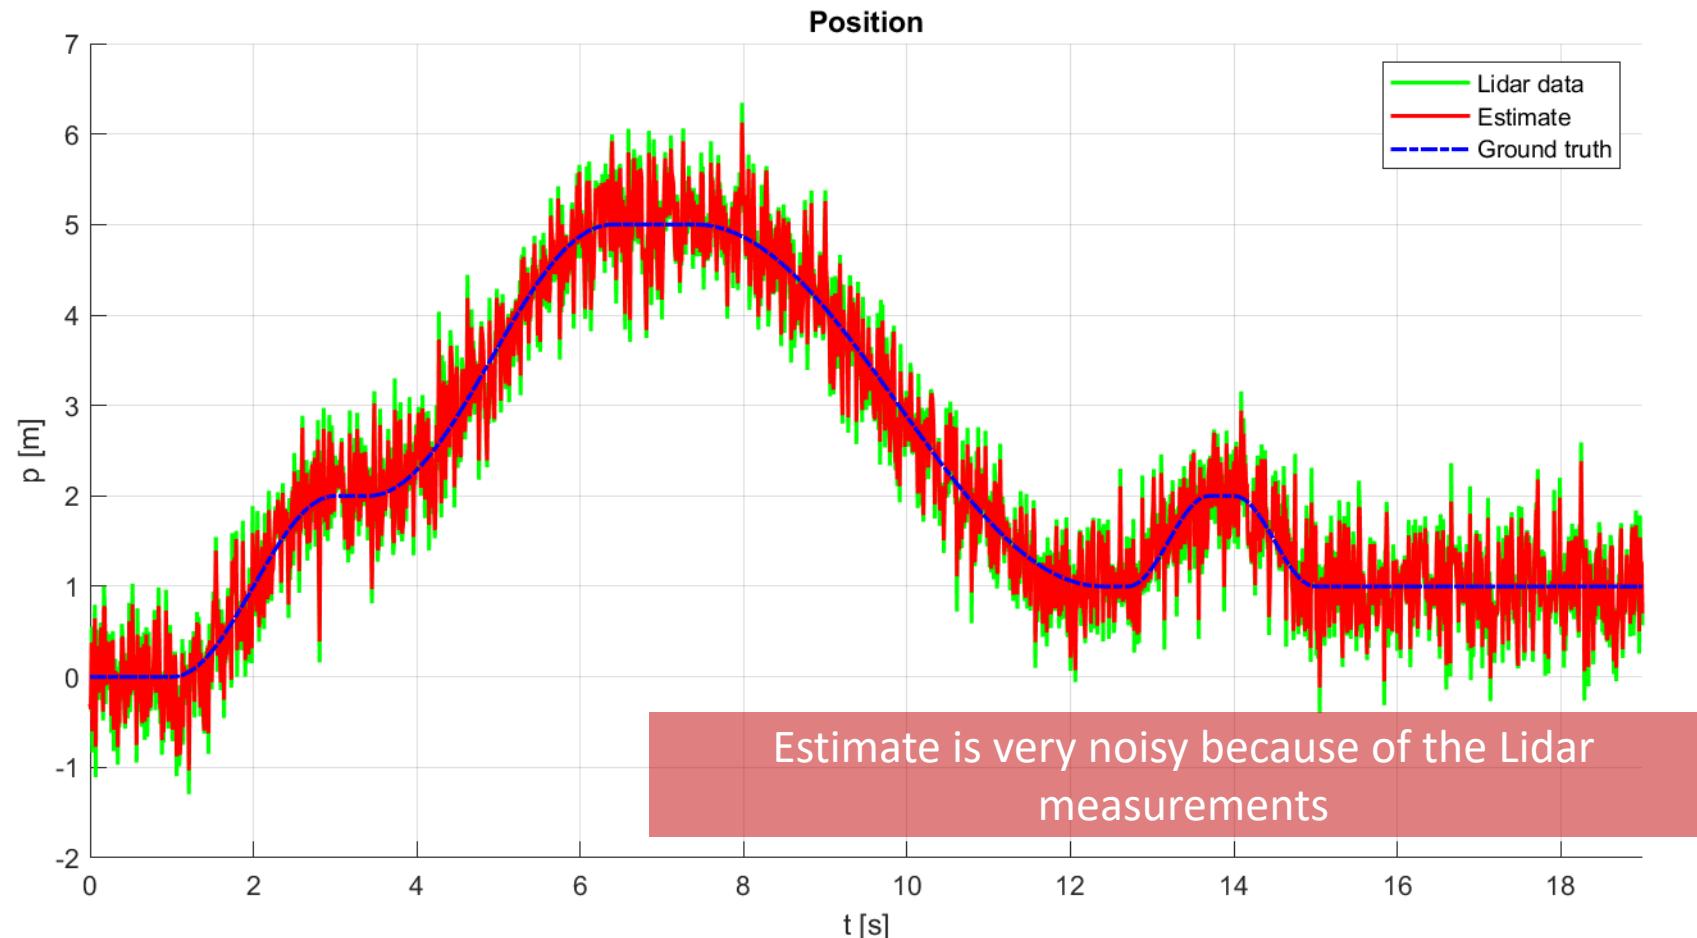
Sensor fusion example: Position estimation



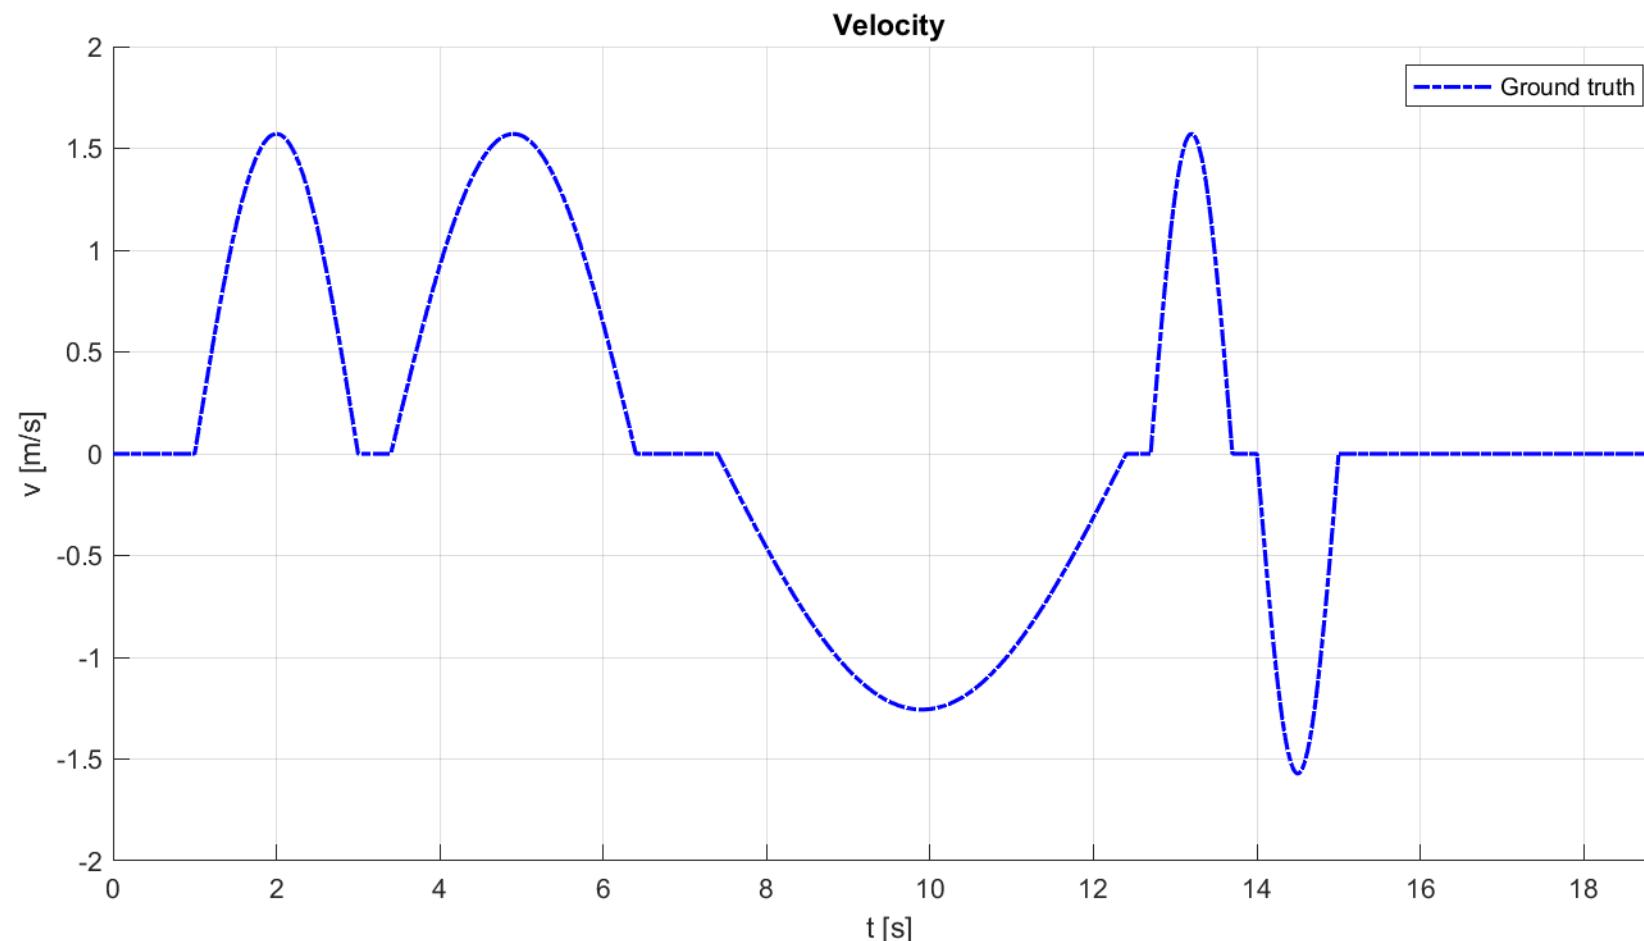
Sensor fusion example: Position estimation



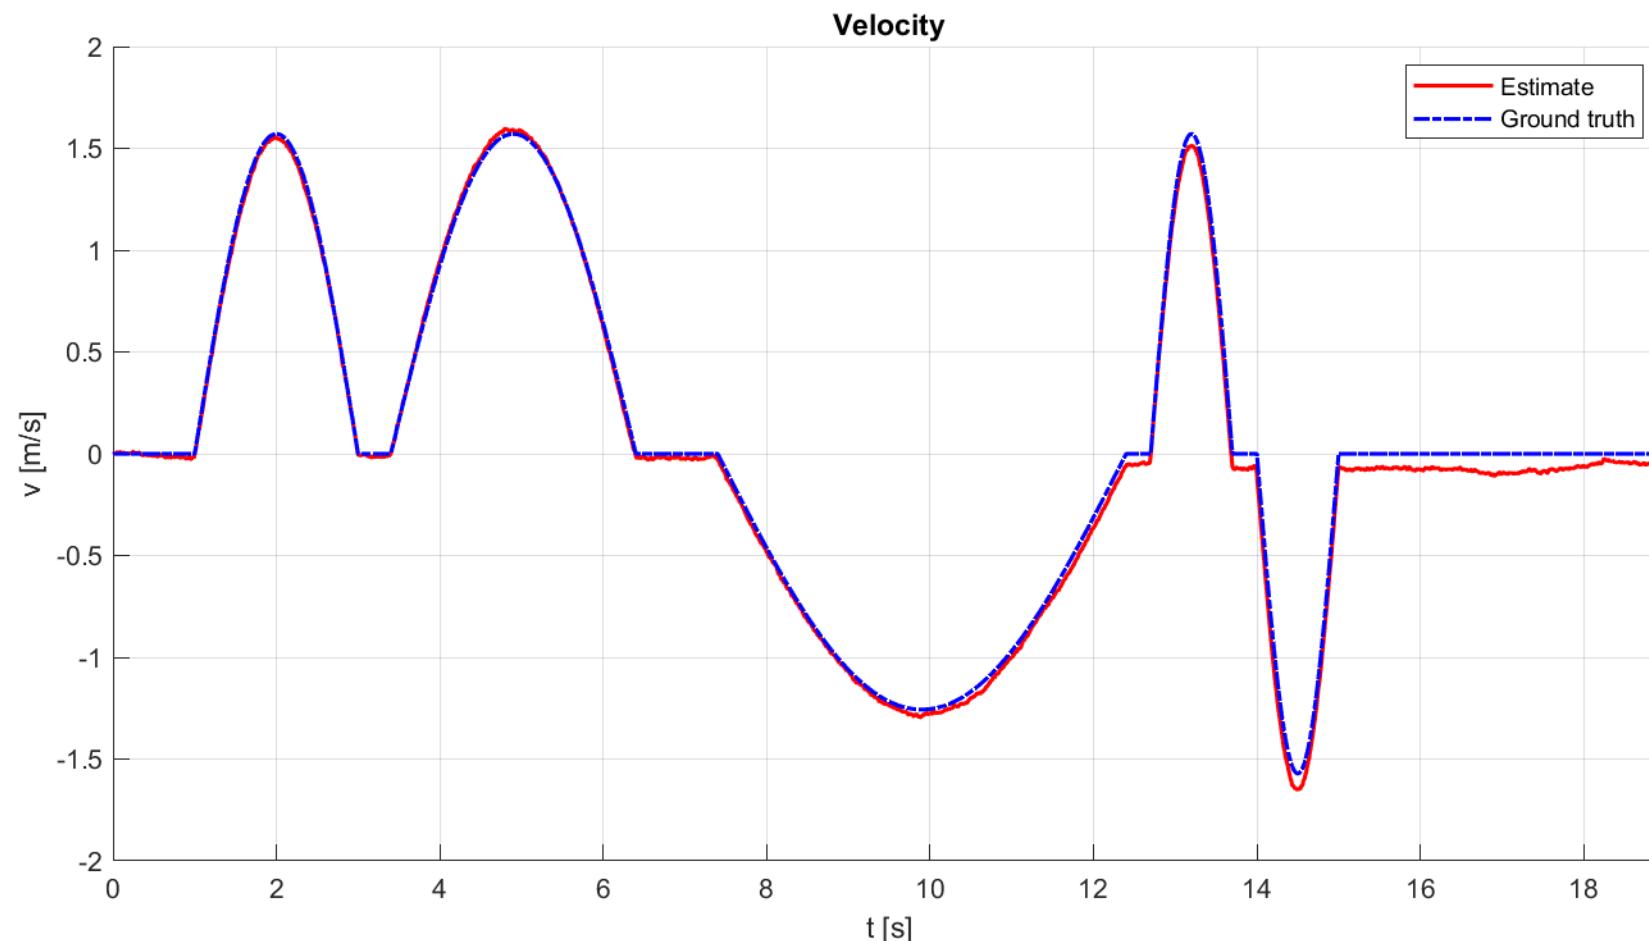
Sensor fusion example: Position estimation



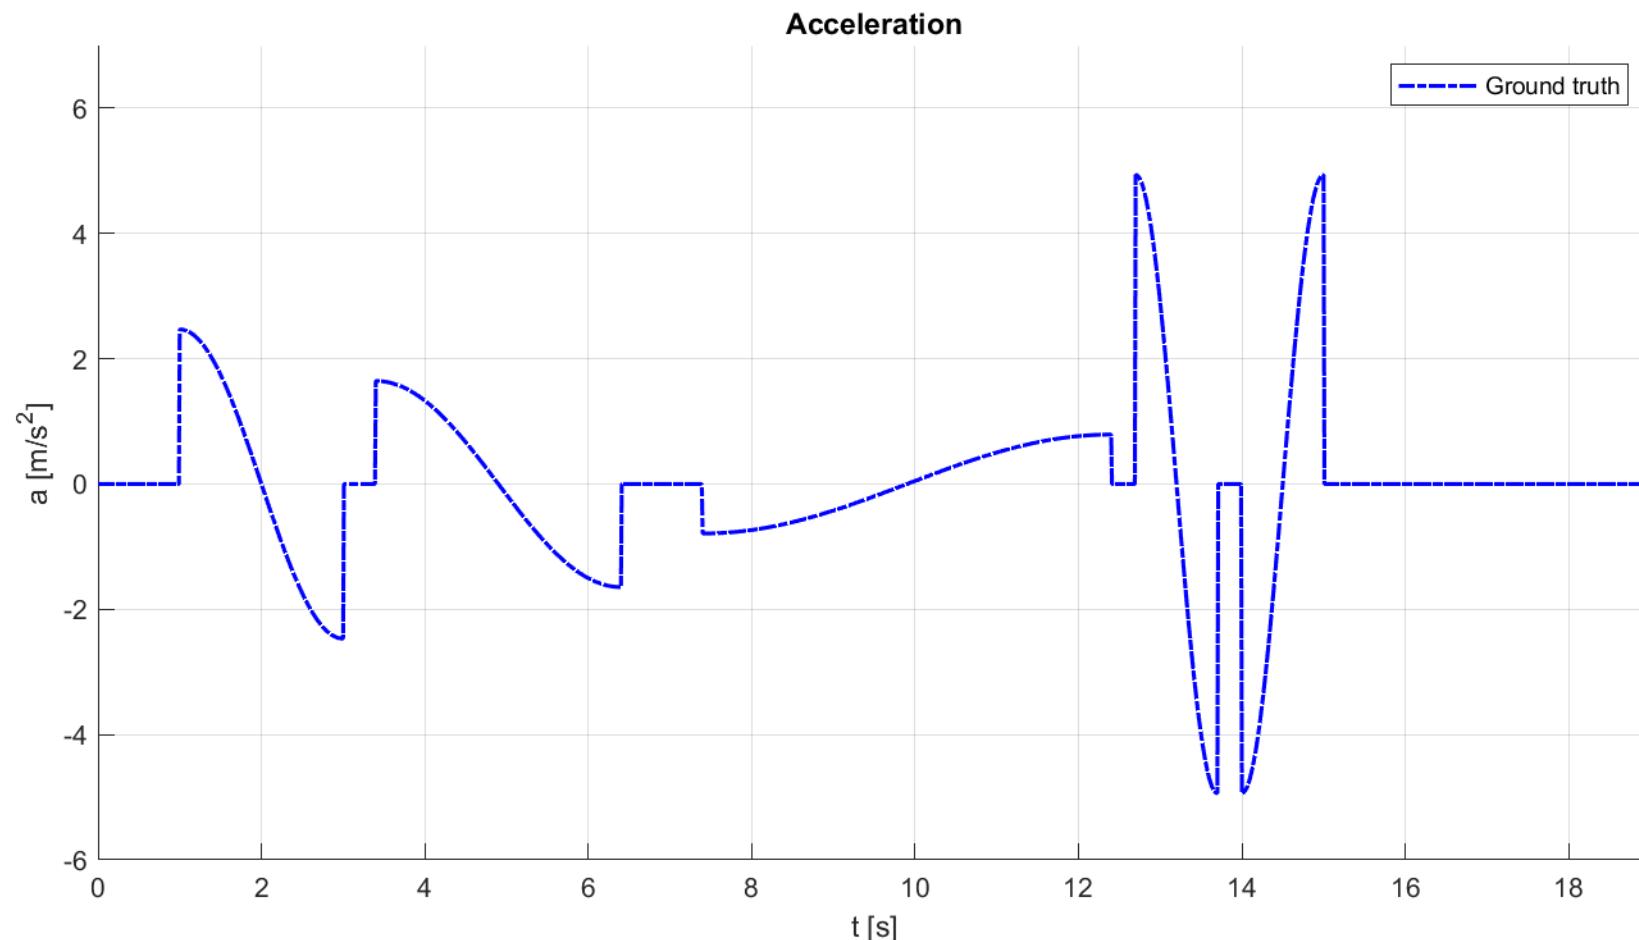
Sensor fusion example: Velocity estimation



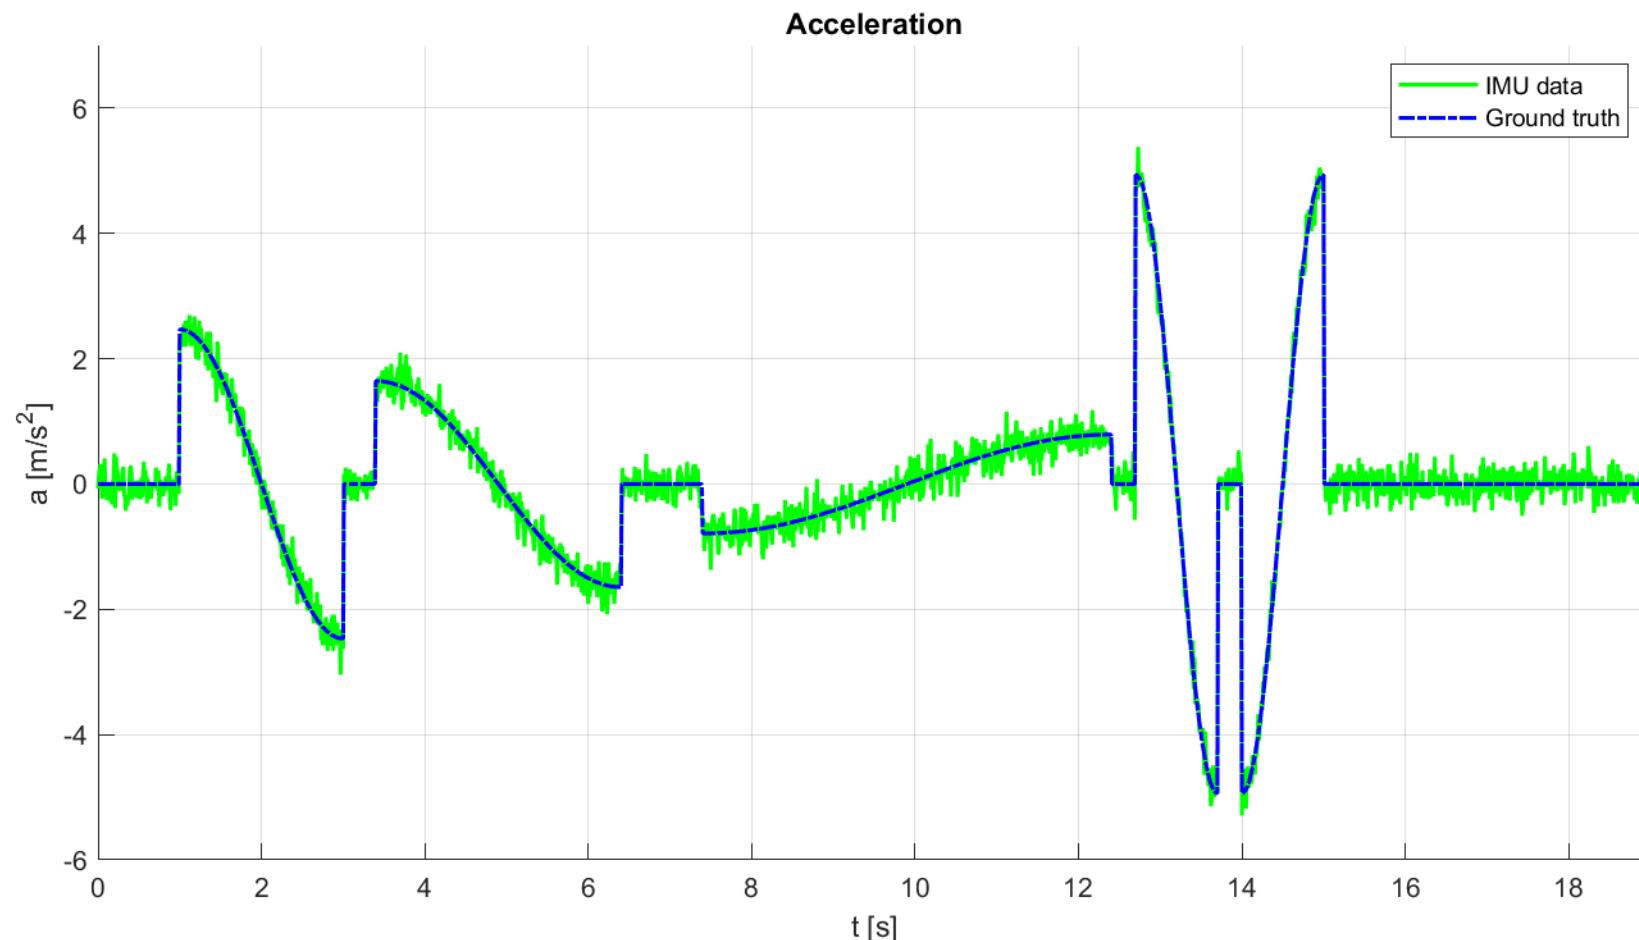
Sensor fusion example: Velocity estimation



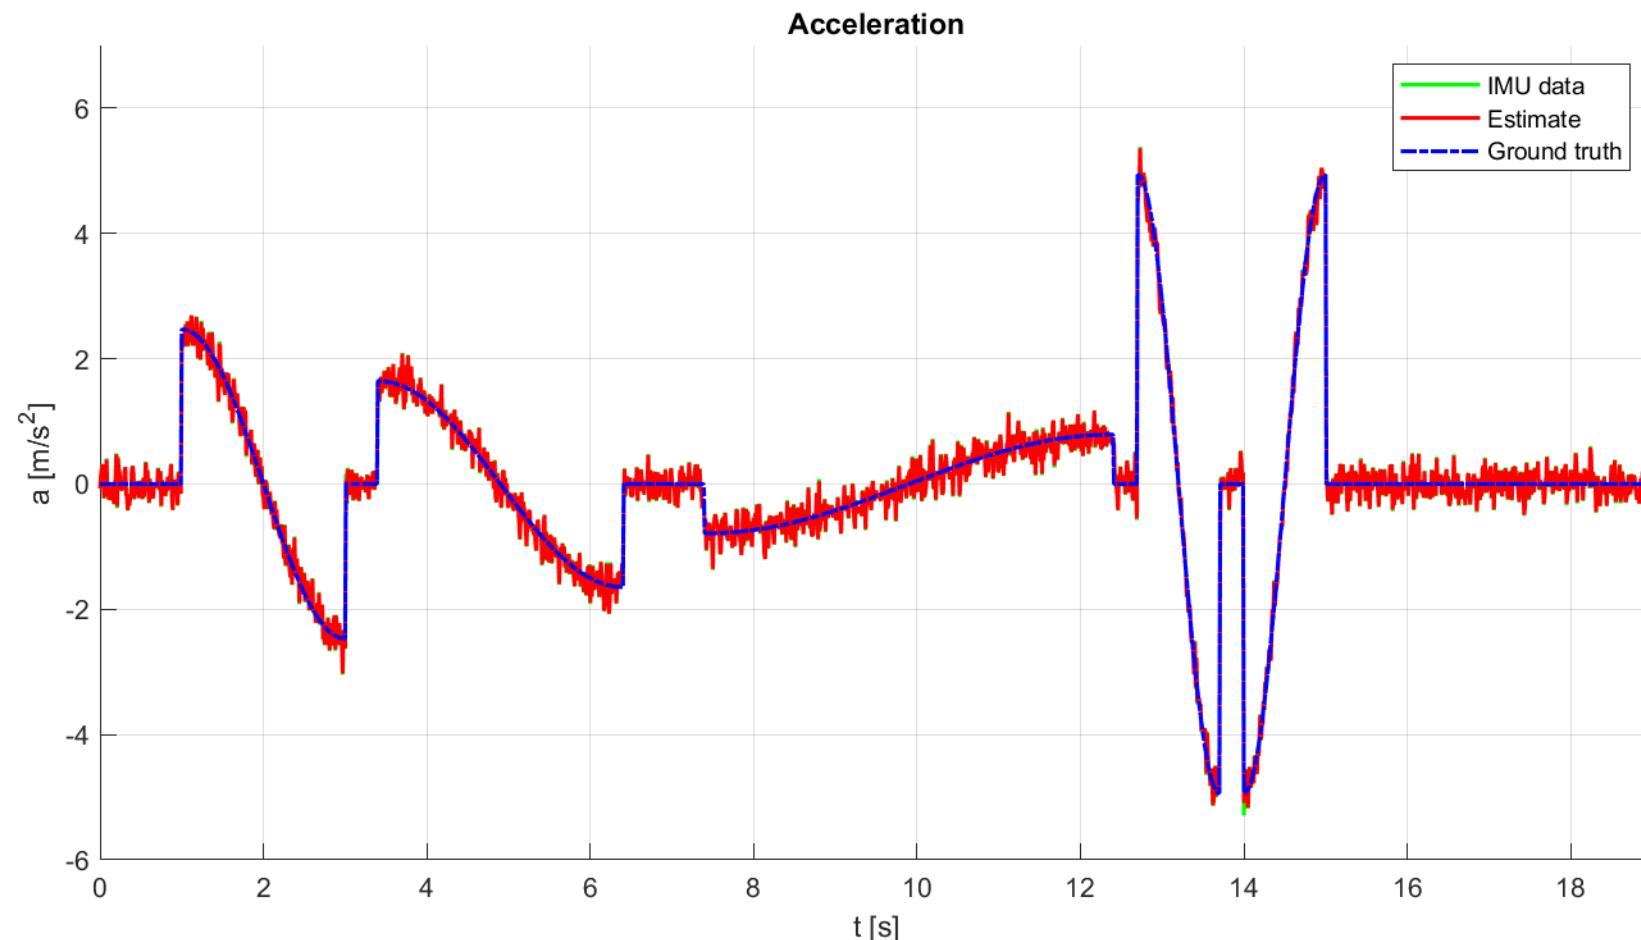
Sensor fusion example: Acceleration estimation



Sensor fusion example: Acceleration estimation



Sensor fusion example: Acceleration estimation



Sensor fusion example: Measurement model

- In previous scenario – the position estimate is quite noisy (because of low precision of the Lidar measurements)
- Therefore, in second scenario, position is measured with Lidar and GPS

$$\begin{bmatrix} p_{lidar} \\ p_{gps} \\ a_{imu} \end{bmatrix}_t = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p \\ v \\ a \end{bmatrix}_t + \begin{bmatrix} \delta_{lidar} \\ \delta_{gps} \\ \delta_{imu} \end{bmatrix}_t$$

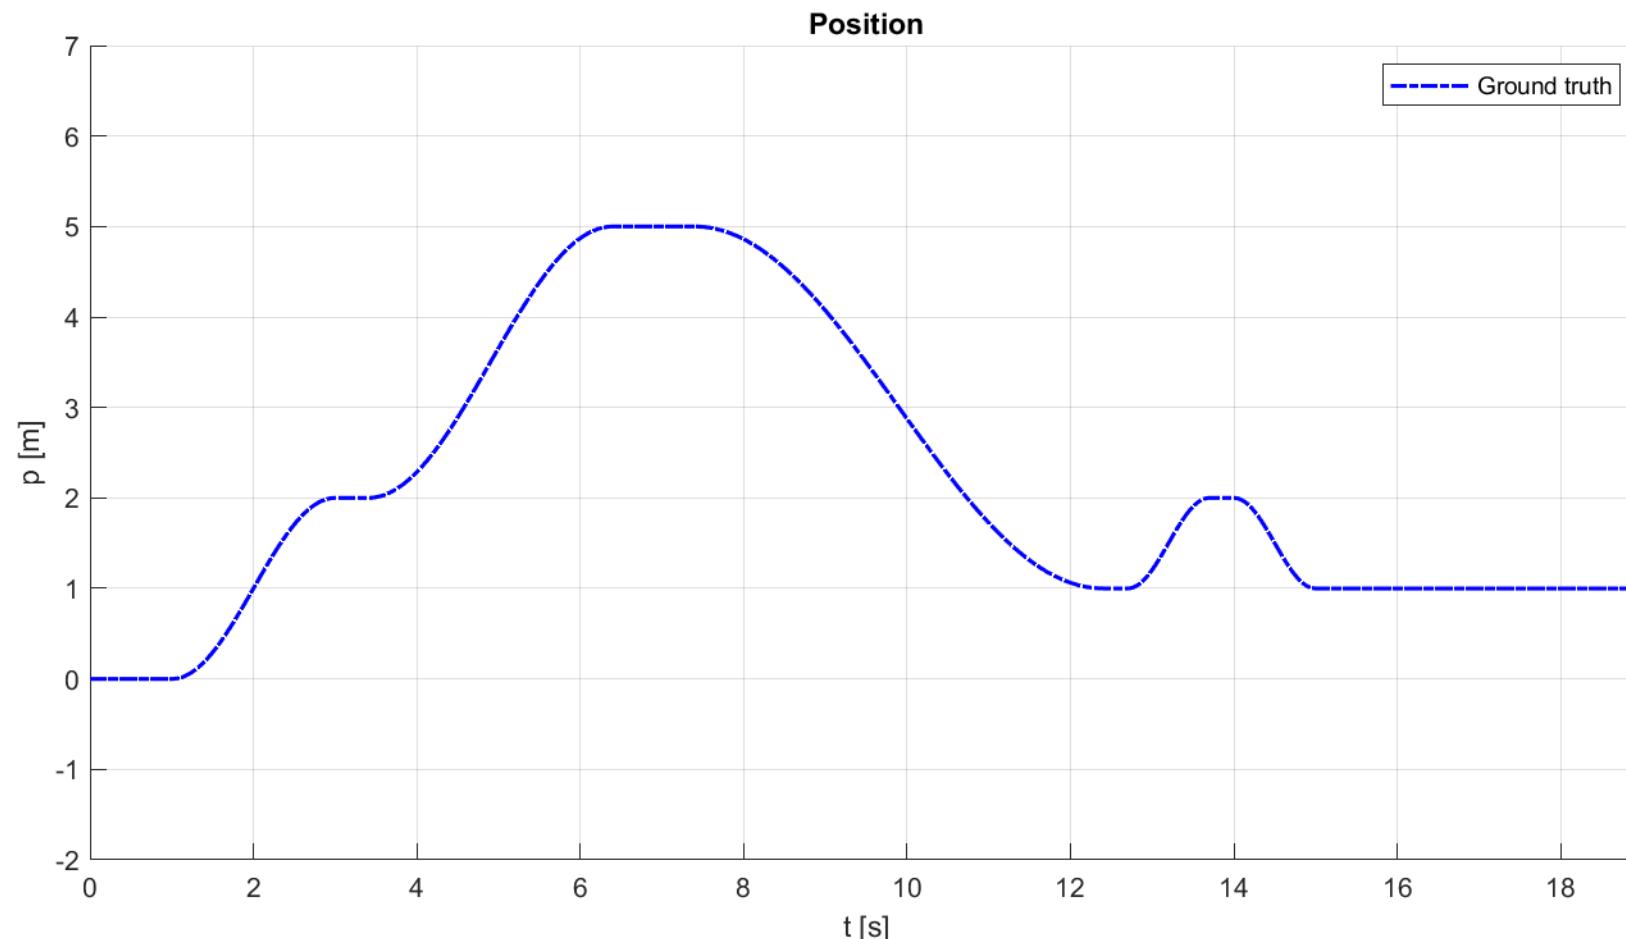
$$z_t = C_t \mu_t + \delta_t$$

Sensor fusion example: Noise model tuning

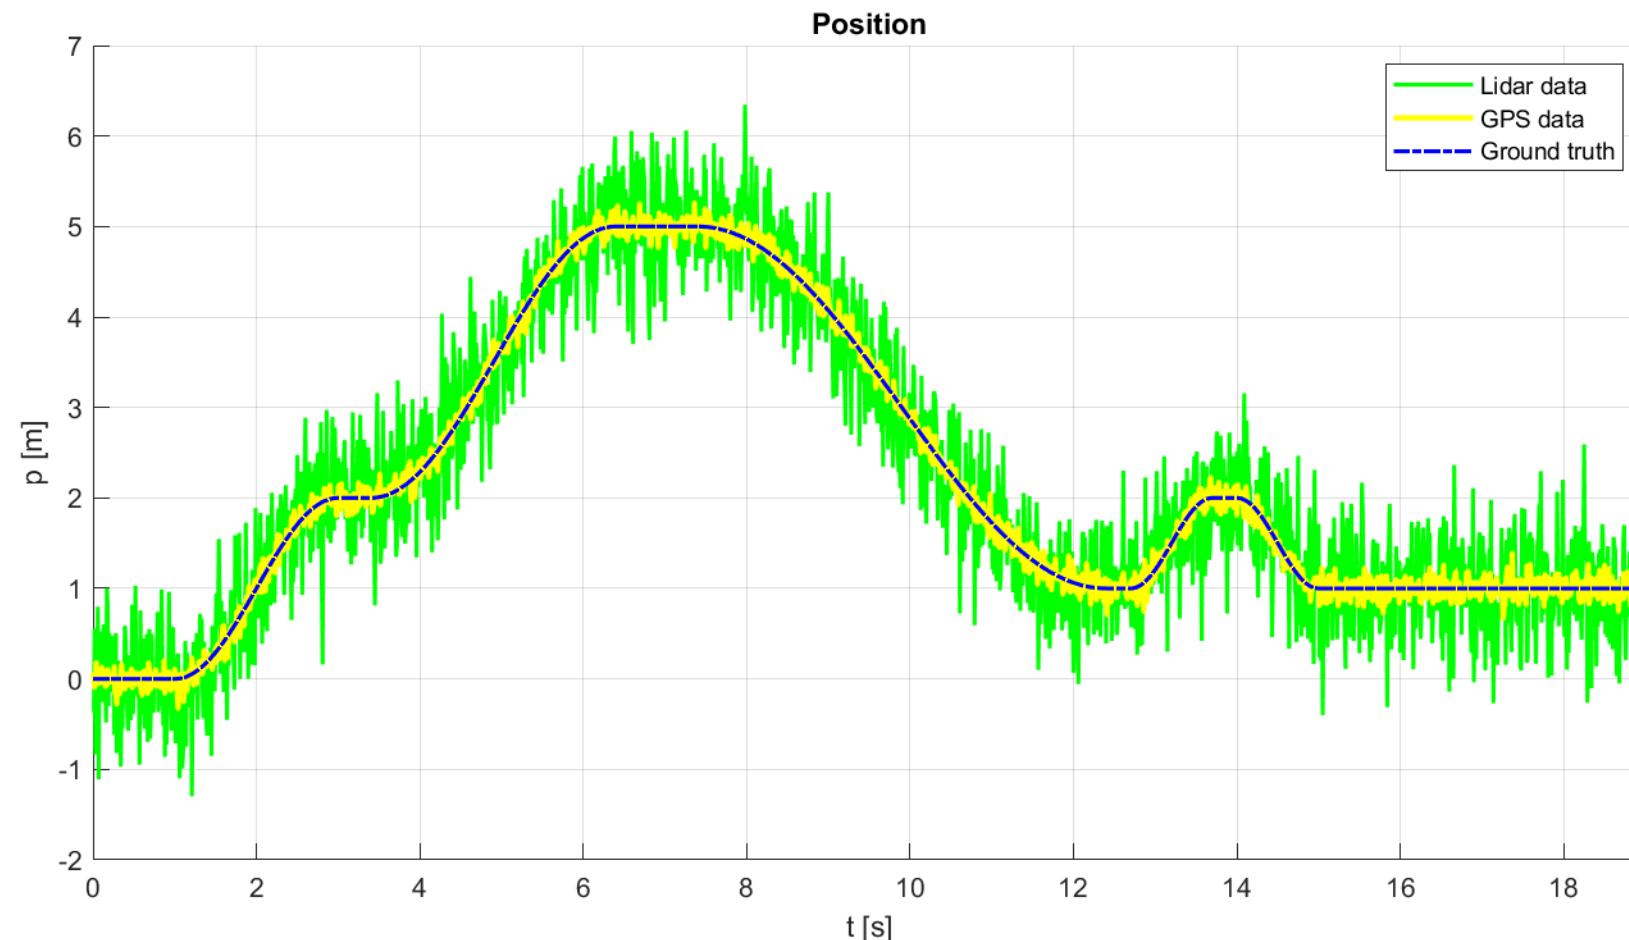
- The measurement noise covariance matrix Q_t for this scenario has additional GPS variance

$$Q_t = \begin{bmatrix} \sigma_{lidar}^2 & 0 & 0 \\ 0 & \sigma_{gps}^2 & 0 \\ 0 & 0 & \sigma_{imu}^2 \end{bmatrix} = \begin{bmatrix} 0.5^2 & 0 & 0 \\ 0 & 0.1^2 & 0 \\ 0 & 0 & 0.2^2 \end{bmatrix}$$

Sensor fusion example: Position estimation



Sensor fusion example: Position estimation



Sensor fusion example: Position estimation



Sensor fusion example: Conclusion

- Problem: Vehicle state estimation using Kalman filter
- The example pointed out:
 - How to create a motion model and a measurement model
 - How to fuse the data from different types of sensors
 - How to set initial state vector and initial covariance matrix
 - How to chose appropriate values for process noise and measurement noise covariance matrices
 - How to achieve more accurate estimation by adding more sensors
 - How fusion of data decreases the overall estimation variance



Useful trick

- Augment the state vector with some auxiliary states and then apply the KF to the augmented state space model
- What can we handle?
 - Colored state noise
 - Colored measurement noise
 - Sensor offset and drifts
 - Sensor faults (sudden offset)
 - Actuator fault (sudden offset)